

Diabetes Dataset Review and ML Model Building

1. Introduction

This notebook performs data preprocessing, model building, and analysis on Dataset #1 (Diabetes Dataset)

2. Importing libraries and loading the Dataset:

```
# Install necessary libraries (Run only once)
!pip install pandas matplotlib seaborn markdown2 weasyprint
```

```
# Import libraries
import pandas as pd # Data manipulation
import matplotlib.pyplot as plt # Data visualization
import seaborn as sns # Statistical data visualization
```

```
Requirement already satisfied: pandas in /usr/local/lib/python3.11/dist-packages (2.2.2)
Requirement already satisfied: matplotlib in /usr/local/lib/python3.11/dist-packages (3.10.0)
Requirement already satisfied: seaborn in /usr/local/lib/python3.11/dist-packages (0.13.2)
Collecting markdown2
  Downloading markdown2-2.5.3-py3-none-any.whl.metadata (2.1 kB)
Collecting weasyprint
  Downloading weasyprint-65.0-py3-none-any.whl.metadata (3.7 kB)
Requirement already satisfied: numpy>=1.23.2 in /usr/local/lib/python3.11/dist-packages (from pandas) (2.0.2)
Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python3.11/dist-packages (from pandas) (2.8.2)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.11/dist-packages (from pandas) (2025.1)
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.11/dist-packages (from pandas) (2025.1)
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.11/dist-packages (from matplotlib) (1.3.1)
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.11/dist-packages (from matplotlib) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.11/dist-packages (from matplotlib) (4.56.0)
Requirement already satisfied: kiwisolver>=1.3.1 in /usr/local/lib/python3.11/dist-packages (from matplotlib) (1.4.8)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.11/dist-packages (from matplotlib) (24.2)
Requirement already satisfied: pillow>=8 in /usr/local/lib/python3.11/dist-packages (from matplotlib) (11.1.0)
Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.11/dist-packages (from matplotlib) (3.2.1)
Collecting pydyf>=0.11.0 (from weasyprint)
  Downloading pydyf-0.11.0-py3-none-any.whl.metadata (2.5 kB)
Requirement already satisfied: cffi>=0.6 in /usr/local/lib/python3.11/dist-packages (from weasyprint) (1.17.1)
Collecting tinyhtml5>=2.0.0b1 (from weasyprint)
  Downloading tinyhtml5-2.0.0-py3-none-any.whl.metadata (2.9 kB)
Requirement already satisfied: tinycss2>=1.4.0 in /usr/local/lib/python3.11/dist-packages (from weasyprint) (1.4.0)
Collecting cssselect2>=0.8.0 (from weasyprint)
  Downloading cssselect2-0.8.0-py3-none-any.whl.metadata (2.9 kB)
Collecting Pyphen>=0.9.1 (from weasyprint)
  Downloading pyphen-0.17.2-py3-none-any.whl.metadata (3.2 kB)
Requirement already satisfied: pycparser in /usr/local/lib/python3.11/dist-packages (from cffi>=0.6->weasyprint) (2.22)
Requirement already satisfied: webencodings in /usr/local/lib/python3.11/dist-packages (from cssselect2>=0.8.0->weasyprint) (0.5.1)
Collecting brotli>=1.0.1 (from fonttools[woff]>=4.0.0->weasyprint)
  Downloading Brotli-1.1.0-cp311-cp311-manylinux_2_17_x86_64.manylinux2014_x86_64.whl.metadata (5.5 kB)
Collecting zopfli>=0.1.4 (from fonttools[woff]>=4.0.0->weasyprint)
  Downloading zopfli-0.2.3.post1-cp311-cp311-manylinux_2_17_x86_64.manylinux2014_x86_64.whl.metadata (2.9 kB)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.11/dist-packages (from python-dateutil>=2.8.2->pandas) (1.17.0)
Downloading markdown2-2.5.3-py3-none-any.whl (48 kB)
48.5/48.5 kB 3.6 MB/s eta 0:00:00
Downloading weasyprint-65.0-py3-none-any.whl (297 kB)
297.9/297.9 kB 21.6 MB/s eta 0:00:00
Downloading cssselect2-0.8.0-py3-none-any.whl (15 kB)
Downloading pydyf-0.11.0-py3-none-any.whl (8.1 kB)
Downloading pyphen-0.17.2-py3-none-any.whl (2.1 MB)
2.1/2.1 MB 55.7 MB/s eta 0:00:00
Downloading tinyhtml5-2.0.0-py3-none-any.whl (39 kB)
Downloading Brotli-1.1.0-cp311-cp311-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (2.9 MB)
2.9/2.9 MB 59.7 MB/s eta 0:00:00
Downloading zopfli-0.2.3.post1-cp311-cp311-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (850 kB)
850.6/850.6 kB 38.1 MB/s eta 0:00:00
Installing collected packages: brotli, zopfli, tinyhtml5, Pyphen, pydyf, markdown2, cssselect2, weasyprint
Successfully installed Pyphen-0.17.2 brotli-1.1.0 cssselect2-0.8.0 markdown2-2.5.3 pydyf-0.11.0 tinyhtml5-2.0.0 weasyprint-65.0 zopfli-0.2.3
```


```
# dataset 1
```

```
import pandas as pd
```


```
# Use the RAW URL, not the repository URL
url = "https://raw.githubusercontent.com/AsmaShaikhTMU/Projects/main/diabetes_prediction_dataset.csv"

# Read the CSV file directly from GitHub
df = pd.read_csv(url)

# Display the first few rows
df.head()
```



	gender	age	hypertension	heart_disease	smoking_history	bmi	HbA1c_level	blood_glucose_level	diabetes
0	Female	80.0	0	1	never	25.19	6.6	140	0
1	Female	54.0	0	0	No Info	27.32	6.6	80	0
2	Male	28.0	0	0	never	27.32	5.7	158	0
3	Female	36.0	0	0	current	23.45	5.0	155	0
4	Male	76.0	1	1	current	20.14	4.8	155	0



Next steps:

[Generate code with df](#)

[View recommended plots](#)

[New interactive sheet](#)


```
# dataset 2

import pandas as pd # Use pd instead of pd2

# Correct RAW URL
url = "https://raw.githubusercontent.com/AsmaShaikhTMU/Projects/main/diabetes_012_health_indicators_BRFSS2015.csv"

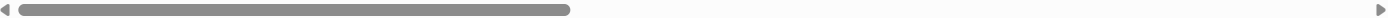
# Read CSV file
df2 = pd.read_csv(url)

# Display first few rows
df2.head()
```



	Diabetes_012	HighBP	HighChol	CholCheck	BMI	Smoker	Stroke	HeartDiseaseorAttack	PhysActivity	Fruits	...	AnyHealthcare	NoDoct
0	0.0	1.0	1.0	1.0	40.0	1.0	0.0	0.0	0.0	0.0	...	1.0	
1	0.0	0.0	0.0	0.0	25.0	1.0	0.0	0.0	1.0	0.0	...	0.0	
2	0.0	1.0	1.0	1.0	28.0	0.0	0.0	0.0	0.0	1.0	...	1.0	
3	0.0	1.0	0.0	1.0	27.0	0.0	0.0	0.0	1.0	1.0	...	1.0	
4	0.0	1.0	1.0	1.0	24.0	0.0	0.0	0.0	1.0	1.0	...	1.0	

5 rows × 22 columns



> 3. Cleaning and Preprocessing the Dataset:

[] ↳ 8 cells hidden

✓ 4. Exploration of Data

✓ I Univariate Analysis :

> Gender

[] ↳ 4 cells hidden

> Age

[] ↳ 2 cells hidden

> Hypertension

[] ↳ 3 cells hidden

> Heart Disease

[] ↳ 3 cells hidden

> Smoke History

[] ↳ 4 cells hidden

> BMI

[] ↳ 4 cells hidden

> HbA1c_level


[] ↳ 5 cells hidden

> Blood Glucose Level

[] ↳ 2 cells hidden

▼ Target variable diabetes

```
#Target variable diabetes
labels = list(df['diabetes'].value_counts().index)
num_var = list(df['diabetes'].value_counts().values)
print(labels)
```

 [0, 1]


```
# Count occurrences of diabetic vs non-diabetic cases
diabetes_counts = df['diabetes'].value_counts().to_dict() # Convert to dictionary to avoid KeyError

# Ensure both labels exist in the correct order
sizes = [diabetes_counts.get(0, 0), diabetes_counts.get(1, 0)] # Handles cases where values might be missing
labels = ['Non-Diabetic', 'Diabetic']
colors = ['#1f77b4', '#ff7f0e'] # Custom colors for clarity
explode = (0, 0.1) # Slightly separate "Diabetic" slice for emphasis

# Create the pie chart
plt.figure(figsize=(6, 6))
plt.pie(sizes, labels=labels, colors=colors, autopct='%1.2f%%', startangle=140, explode=explode)
plt.title("Percentage of Diabetic vs Non-Diabetic Patients")
plt.legend(labels, loc="upper right")
plt.show()
```



Percentage of Diabetic vs Non-Diabetic Patients



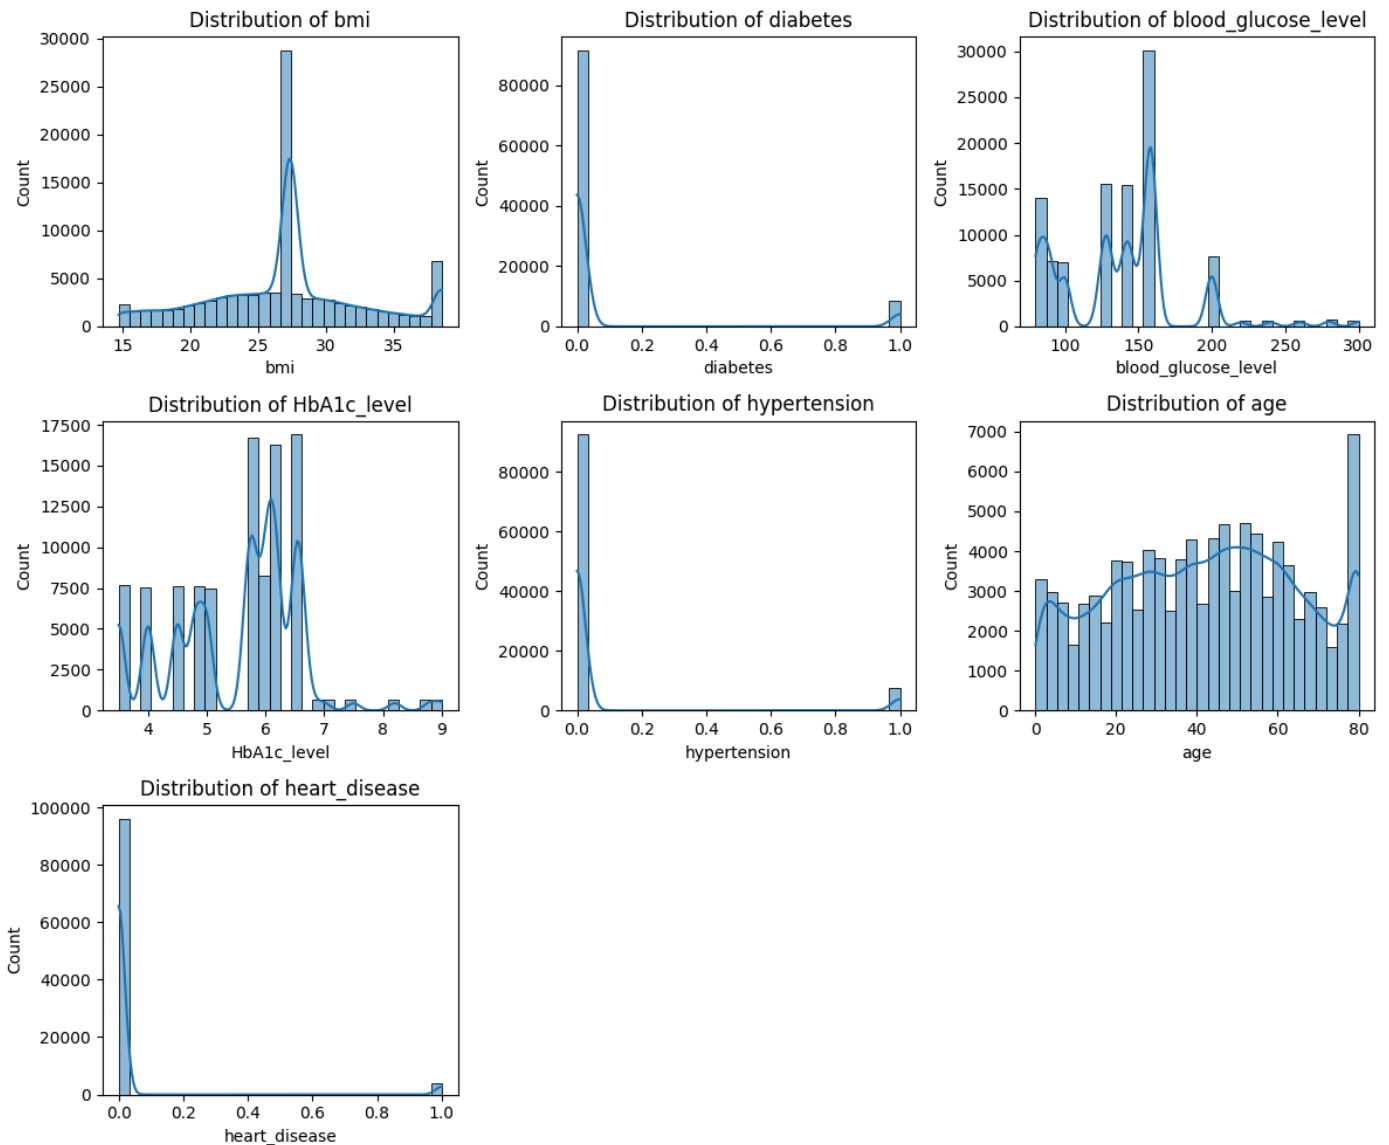
**Given that diabetic cases about 8% of the dataset (high imbalance), the best approach should enhance the minority class representation.

Method 1: oversampling with SMOTE.

Method 2: Class Weight Adjustment in Models helps the model give more importance to the minority class.

Method 3: Use Recall & F1 Score for Model Evaluation (recall to ensure diabetic cases are detected and F1-score to balance precision & recall (instead of accuracy)**

```
# Distribution plots for all variables
variables = ['bmi', 'diabetes', 'blood_glucose_level', 'HbA1c_level', 'hypertension', 'age', 'heart_disease']
plt.figure(figsize=(12, 10))
for i, var in enumerate(variables, 1):
    plt.subplot(3, 3, i)
    sns.histplot(df[var], kde=True, bins=30)
    plt.title(f"Distribution of {var}")
plt.tight_layout()
plt.show()
```



Blood Glucose is multi-modal distribution.

BMI is right-skewed distribution.

HbA1c Level is uneven distribution.


The age distribution appears relatively uniform but with peaks around 60-80 years. A higher count of older individuals suggests a dataset focused on age-related health conditions.

There is a notable imbalance, with far fewer cases of heart disease.

✓ II. Bivariate Analysis

In this section, the relationships between two variables and if possible, create new features combining them for better visualization

```
#Correlation heatmap
plt.figure(figsize=(8, 6))
sns.heatmap(df.corr(numeric_only=True), annot=True, cmap='coolwarm', fmt=".2f", linewidths=0.5)
plt.title("Feature Correlation Heatmap")
plt.show()
```



Diabetes Correlation:

Blood Glucose Level (0.42): Strongest correlation - diabetes is diagnosed based on high blood sugar.

HbA1c Level (0.40): Strong correlation - HbA1c reflects long-term blood glucose levels.

Age (0.26): Moderate correlations

Hypertension (0.20): Moderate correlations

BMI (0.21): Moderate correlations

✓ 1. HbA1c_level vs. diabetes:

Question: What percentage of individuals with HbA1c levels above 6.5% have diabetes?

HbA1c_level

A hemoglobin A1C (HbA1C) test is a blood test that shows what your average blood sugar (glucose) level was over the past two to three months. we will create a new feature based on the value of (HbA1C)

HbA1c level initial diagnosis

1. < 5.7 Normal
2. 5.7 – 6.4 Prediabetes
3. greater or equal 6.5 Diabetes

Resource : <https://www.cdc.gov/diabetes/managing/managing-blood-sugar/a1c.html>

```
# Calculate percentage of individuals with HbA1c > 6.5% who have diabetes
```

```
diabetic_hba1c = df[(df['HbA1c_level'] > 6.5) & (df['diabetes'] == 1)].shape[0]
total_hba1c_above_6_5 = df[df['HbA1c_level'] > 6.5].shape[0]
```

```
if total_hba1c_above_6_5 > 0:
    percentage_diabetic_hba1c = (diabetic_hba1c / total_hba1c_above_6_5) * 100
    print(f"Percentage of individuals with HbA1c > 6.5% who have diabetes: {percentage_diabetic_hba1c:.2f}%")
```

```
else:
    print("No individuals with HbA1c > 6.5% in the dataset.")

↵ Percentage of individuals with HbA1c > 6.5% who have diabetes: 36.82%
```

The dataset **HbA1c_level** value increases, number of records with diabetes increases ## specially when it >= 6.5

```
# Define HbA1c level categories based on diabetes classification
hba1c_bins = [0, 5.7, 6.4, float('inf')]
hba1c_labels = ['Normal', 'Prediabetes', 'Diabetes']
df['HbA1c_Category'] = pd.cut(df['HbA1c_level'], bins=hba1c_bins, labels=hba1c_labels)

# Calculate percentage of diabetics in each HbA1c category
hba1c_diabetes_counts = df.groupby('HbA1c_Category', observed=True)['diabetes'].mean() * 100


# Create a structured table
hba1c_table = pd.DataFrame({
    'HbA1c Level Category': hba1c_labels,
    'Prediction': [f"{hba1c_diabetes_counts.get(category, 0):.2f}% have diabetes" for category in hba1c_labels]
})

# Display the table
print(hba1c_table)
```

```
↵
```

HbA1c Level Category	Prediction
0 Normal	1.52% have diabetes
1 Prediabetes	7.91% have diabetes
2 Diabetes	24.96% have diabetes

```
plt.figure(figsize=(8, 6))
sns.boxplot(x="diabetes", y="HbA1c_level", data=df, hue="diabetes", palette={0: "#1f77b4", 1: "#ff7f0e"}, legend=False)
plt.xticks([0, 1], ['Non-Diabetic', 'Diabetic'])
plt.title("HbA1c Level Distribution in Diabetic vs. Non-Diabetic Patients")
plt.xlabel("Diabetes Status")
plt.ylabel("HbA1c Level")
plt.show()
```



> 2. bmi vs diabetes:

[] ↵ 7 cells hidden

3. blood_glucose_level vs diabetes

[] ↳ 17 cells hidden

4. hypertension vs bmi

[] ↳ 4 cells hidden

5. Age vs Heart Disease

[] ↳ 3 cells hidden

✓ Hypothesis Testing

✓ T-test: Compare blood glucose levels between diabetic and non-diabetic individuals.

Use a T-Test when: Target (diabetic vs. non-diabetic patients). Continuous variable (e.g., blood glucose levels, BMI, HbA1c levels).

```
import scipy.stats as stats

glucose_diabetic = df[df['diabetes'] == 1]['blood_glucose_level'] # Use actual column names
glucose_non_diabetic = df[df['diabetes'] == 0]['blood_glucose_level'] # Use actual column names

t_stat, p_value_ttest = stats.ttest_ind(glucose_diabetic, glucose_non_diabetic, equal_var=False)

print(f"T-Test Results (Blood Glucose vs Diabetes):")
print(f"T-Statistic: {t_stat:.4f}, P-Value: {p_value_ttest:.4f}")
print("Significant Difference" if p_value_ttest < 0.05 else "No Significant Difference")
print("-" * 50)
```

```
➡ T-Test Results (Blood Glucose vs Diabetes):
T-Statistic: 94.7949, P-Value: 0.0000
Significant Difference
-----
```

Since the p-value is much lower than 0.05, we reject the null hypothesis. There is a significant difference in blood glucose levels between diabetic and non-diabetic individuals. Blood glucose is a strong differentiator for diabetes.

✓ Chi-Square Test: Check if hypertension and diabetes are significantly associated.

The Chi-Square Test is a statistical test used to determine whether there is a significant association between two categorical variables

```
contingency_table = pd.crosstab(df['hypertension'], df['diabetes'])
chi2_stat, p_value_chi2, dof, expected = stats.chi2_contingency(contingency_table)

print(f"Chi-Square Test Results (Hypertension vs Diabetes):")
print(f"Chi-Square Statistic: {chi2_stat:.4f}, P-Value: {p_value_chi2:.4f}")
print("Significant Association" if p_value_chi2 < 0.05 else "No Significant Association")
print("-" * 50)
```

```
➡ Chi-Square Test Results (Hypertension vs Diabetes):
Chi-Square Statistic: 3910.7085, P-Value: 0.0000
Significant Association
-----
```

There is strong statistical association between hypertension and diabetes. This means that individuals with hypertension are significantly more likely to have diabetes.

Predictive Modeling

✓ 5. Train-test split and standardized scaler

```
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
import pandas as pd

# Select only numerical features for scaling
numerical_features = df.select_dtypes(include=['number']).columns
X = df[numerical_features].drop('diabetes', axis=1)
y = df['diabetes']

# Split data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Apply StandardScaler only to numerical features
scale = StandardScaler()
X_train = scale.fit_transform(X_train)
X_test = scale.transform(X_test)

print("X_train:", X_train.shape, "\ny_train:", y_train.shape)
print("X_test:", X_test.shape, "\ny_test:", y_test.shape)
```

➞ X_train: (80000, 6)
y_train: (80000,)
X_test: (20000, 6)
y_test: (20000,)

✓ Logistic Regression Classifier

```
## Designing Logistic Regression Classifier
import numpy as np # Import numpy and assign it to the alias 'np'
from sklearn.linear_model import LogisticRegression # Import the LogisticRegression class
from sklearn.metrics import accuracy_score, mean_absolute_error, mean_squared_error, precision_score, recall_score, f1_score # Import necess

LogReg = LogisticRegression()
LogReg.fit(X_train, y_train)
LogReg_pred = LogReg.predict(X_test)
LogReg_acc = accuracy_score(y_test, LogReg_pred)
LogReg_mae = mean_absolute_error(y_test, LogReg_pred)
LogReg_mse = mean_squared_error(y_test, LogReg_pred)
LogReg_rmse = np.sqrt(mean_squared_error(y_test, LogReg_pred))
LogReg_precision = precision_score(y_test, LogReg_pred)
LogReg_recall = recall_score(y_test, LogReg_pred)
LogReg_f1 = f1_score(y_test, LogReg_pred)

## Designing Logistic Regression Classifier
LogReg = LogisticRegression()
LogReg.fit(X_train, y_train)
LogReg_pred = LogReg.predict(X_test)
LogReg_acc = accuracy_score(y_test, LogReg_pred)
LogReg_mae = mean_absolute_error(y_test, LogReg_pred)
LogReg_mse = mean_squared_error(y_test, LogReg_pred)
LogReg_rmse = np.sqrt(mean_squared_error(y_test, LogReg_pred))
LogReg_precision = precision_score(y_test, LogReg_pred)
LogReg_recall = recall_score(y_test, LogReg_pred)
LogReg_f1 = f1_score(y_test, LogReg_pred)

## Printing the results
from sklearn.metrics import classification_report # Import classification_report

print("The accuracy for Logistic Regression is", LogReg_acc)
print("The classification report using Logistic Regression is:")
print(classification_report(y_test, LogReg_pred))
```

➞ The accuracy for Logistic Regression is 0.95875
The classification report using Logistic Regression is:

	precision	recall	f1-score	support
0	0.96	0.99	0.98	18292
1	0.86	0.61	0.72	1708

accuracy			0.96	20000
macro avg	0.91	0.80	0.85	20000
weighted avg	0.96	0.96	0.96	20000

Observation: accuracy alone is not always the best metric when dealing with imbalanced datasets

Precision: When predicting "Diabetes," it's correct 86% of the time.


Recalls: The model correctly identifies only 61% of diabetics.

F1-Score (Balanced Precision & Recall)

- No Diabetes (0): 0.98 (Excellent balance).
- Diabetes (1): 0.72 (Moderate but could be improved).

```
## Confusion Matrix
from sklearn.metrics import confusion_matrix
LogReg_cm = confusion_matrix(y_test, LogReg_pred)
sns.heatmap(LogReg_cm/np.sum(LogReg_cm), annot = True, fmt = '.2%', cmap = 'Blues')
plt.title("Logisitic Regression Confusion Matrix")
```

Text(0.5, 1.0, 'Logisitic Regression Confusion Matrix')



Double-click (or enter) to edit

Actual No Diabetes (0): ✓ 90.64% (True Negatives) → Correctly predicted as non-diabetic. ✗ 0.83% (False Positives) → Incorrectly predicted as diabetic.

Actual Diabetes (1): ✗ 3.30% (False Negatives) → Incorrectly predicted as non-diabetic. ✓ 5.24% (True Positives) → Correctly predicted as diabetic.

✓ K-Nearest Neighbour Regression Classifier

```
## Designing KNN Classifier
from sklearn.neighbors import KNeighborsClassifier # Import KNeighborsClassifier

KNN = KNeighborsClassifier()
KNN.fit(X_train, y_train)
KNN_pred = KNN.predict(X_test)
KNN_acc = accuracy_score(y_test, KNN_pred)
KNN_mae = mean_absolute_error(y_test, KNN_pred)
KNN_mse = mean_squared_error(y_test, KNN_pred)
KNN_rmse = np.sqrt(mean_squared_error(y_test, KNN_pred))
KNN_precision = precision_score(y_test, KNN_pred)
KNN_recall = recall_score(y_test, KNN_pred)
KNN_f1 = f1_score(y_test, KNN_pred)

from sklearn.metrics import accuracy_score, mean_absolute_error, mean_squared_error, precision_score, recall_score, f1_score, classification_report

##Printing the results
print("The accuracy for KNeighbors is", KNN_acc)
print("The classification report using KNeighbors is:" ),
print(classification_report(y_test, KNN_pred))
```

```
→ The accuracy for KNeighbors is 0.96625
The classification report using KNeighbors is:
```

	precision	recall	f1-score	support
0	0.97	0.99	0.98	18292
1	0.92	0.66	0.77	1708
accuracy			0.97	20000
macro avg	0.95	0.83	0.88	20000
weighted avg	0.97	0.97	0.96	20000

F1-Score:

No Diabetes (0): 0.98 (Excellent performance).

Diabetes (1): 0.77 (Better than Logistic Regression but still room for improvement).

✓ Decision Tree Classifier

```
##Designing Decision Tree Classifier
from sklearn.tree import DecisionTreeClassifier # Import DecisionTreeClassifier
DecTree = DecisionTreeClassifier()
DecTree.fit(X_train, y_train)
DecTree_pred = DecTree.predict(X_test)
DecTree_acc = accuracy_score(y_test, DecTree_pred)
DecTree_precision = precision_score(y_test, DecTree_pred)
DecTree_recall = recall_score(y_test, DecTree_pred)
DecTree_f1 = f1_score(y_test, DecTree_pred)
```

```
##Printing the results
print("The accuracy for Decision Tree is", DecTree_acc)
print("The classification report using Decision Tree is:")
print(classification_report(y_test, DecTree_pred))
```

```
→ The accuracy for Decision Tree is 0.9557
The classification report using Decision Tree is:
```

	precision	recall	f1-score	support
0	0.98	0.98	0.98	18292
1	0.74	0.74	0.74	1708
accuracy			0.96	20000
macro avg	0.86	0.86	0.86	20000
weighted avg	0.96	0.96	0.96	20000

Better recall for diabetes compared to Logistic Regression (61%) and KNN (66%), meaning the model is less likely to miss actual diabetics.

F1-Score (Harmonic mean of Precision & Recall)

- No Diabetes (0): 0.98 (Excellent performance).
- Diabetes (1): 0.74 (Best so far compared to other models).

Balanced performance in identifying diabetes and non-diabetes cases.

✓ Random Forest Classifier

```
##Designing Random Forest Classifier
from sklearn.ensemble import RandomForestClassifier # Import RandomForestClassifier
RFTree = RandomForestClassifier()
RFTree.fit(X_train, y_train)
RFTree_pred = RFTree.predict(X_test)
RFTree_acc = accuracy_score(y_test, RFTree_pred)
RFTree_precision = precision_score(y_test, RFTree_pred)
RFTree_recall = recall_score(y_test, RFTree_pred)
RFTree_f1 = f1_score(y_test, RFTree_pred)
```

```
##Printing the results
print("The accuracy for Random Forest is", RFTree_acc)
print("The classification report using Random Forest is:")
print(classification_report(y_test, RFTree_pred))
```

```
→ The accuracy for Random Forest is 0.96925
The classification report using Random Forest is:
```

	precision	recall	f1-score	support
0	0.97	0.99	0.98	18292
1	0.92	0.70	0.80	1708
accuracy			0.97	20000
macro avg	0.95	0.85	0.89	20000
weighted avg	0.97	0.97	0.97	20000

Recall: Improved recall for diabetics compared to Logistic Regression (61%) and KNN (66%), but slightly lower than Decision Tree (74%).

F1-Score (Harmonic mean of Precision & Recall)

- No Diabetes (0): 0.98 (Excellent performance).
- Diabetes (1): 0.79 (Best so far compared to other models).

Diabetes recall is still lower, meaning the model still misses some diabetics.

✓ Support Vector Machine Classifier

```
#Designing SVM Classifier
from sklearn.svm import SVC
SVM = SVC()
SVM.fit(X_train, y_train)
SVM_pred = SVM.predict(X_test)
SVM_acc = accuracy_score(y_test, SVM_pred)
SVM_precision = precision_score(y_test, SVM_pred)
SVM_recall = recall_score(y_test, SVM_pred)
SVM_f1 = f1_score(y_test, SVM_pred)
```

```
print("The accuracy for SVM is", SVM_acc)
print("The classification report using SVM is:", SVM_acc)
print(classification_report(y_test, SVM_pred))
```

```
→ The accuracy for SVM is 0.96335
The classification report using SVM is: 0.96335
```

	precision	recall	f1-score	support
0	0.96	1.00	0.98	18292
1	0.98	0.58	0.73	1708
accuracy			0.96	20000
macro avg	0.97	0.79	0.86	20000
weighted avg	0.96	0.96	0.96	20000

Recall: 42% of actual diabetic cases were misclassified as non-diabetic (false negatives).

Diabetes recall is much lower than Random Forest (70%) and Decision Tree (74%).

✓ Comparison of Different Models

```
models = pd.DataFrame({
    'Model': ['Logistic Regression', 'KNN Regression', 'Decision Tree', 'Random Forest', 'Support Vector'],
    'Accuracy': [LogReg_acc, KNN_acc, DecTree_acc, RFTree_acc, SVM_acc],
    'Precision': [LogReg_precision, KNN_precision, DecTree_precision, RFTree_precision, SVM_precision],
    'Recall': [LogReg_recall, KNN_recall, DecTree_recall, RFTree_recall, SVM_recall],
    'F1 Score': [LogReg_f1, KNN_f1, DecTree_f1, RFTree_f1, SVM_f1]
})
```

models

	Model	Accuracy	Precision	Recall	F1 Score	
0	Logistic Regression	0.95875	0.863974	0.613583	0.717562	
1	KNN Regression	0.96625	0.922322	0.660422	0.769703	
2	Decision Tree	0.95570	0.738676	0.744731	0.741691	
3	Random Forest	0.96925	0.918774	0.701991	0.795885	
4	Support Vector	0.96335	0.978410	0.583724	0.731206	

Next steps:

[Generate code with models](#)

[View recommended plots](#)

[New interactive sheet](#)

Random Forest: Best F1-score (0.7923), making it the strongest performer

KNN: High precision (92.2%), meaning it correctly identifies diabetics when predicted.

Logistic Regression is the Weakest in Precision & Recall Tradeoff


```
# Create figure
fig = plt.figure(figsize=(15, 9))

# Fixing the warning: Added hue='Model' and legend=False
ax = sns.barplot(data=models,
                 y='Model',
                 x='F1 Score',
                 hue='Model',          # Assigning 'hue' to fix warning
                 palette='Blues_r',
                 legend=False)         # Legend is redundant here

# Title and labels
ax.figure.suptitle('Performance of Models', y=0.91, size=16, color='black', weight='bold')
plt.xlabel('F1 Score')
plt.ylabel('Models')

# Add a reference vertical line for the threshold
plt.axvline(x=0.96, ymin=0, ymax=1,
            linewidth=3, linestyle=":",
            color='#cf7849')

plt.show()
```



✓ CROSS VALIDATION

K-Fold Cross-Validation is used to ensure a model's performance is reliable and not dependent on a single training/test split. Since F1-score is the most balanced metric for imbalanced datasets (like diabetes detection), by apply K-Fold Cross-Validation to the top 3 models with the highest F1-scores

##Performing K-Fold cross Validation for 3 models performed better in F1-Score

```
from sklearn.model_selection import KFold
from sklearn.model_selection import cross_val_score
from statistics import mean, stdev
```

```
cv = KFold(n_splits=10, random_state=1, shuffle=True)
```

```
KNN_scores = cross_val_score(KNN, X, y, scoring='accuracy', cv=cv, n_jobs=-1)
DecTree_scores = cross_val_score(DecTree, X, y, scoring='accuracy', cv=cv, n_jobs=-1)
RFTree_scores = cross_val_score(RFTree, X, y, scoring='accuracy', cv=cv, n_jobs=-1)
```

```
print('Accuracy of CV - KNN: %.4f (%.4f)' % (mean(KNN_scores), stdev(KNN_scores)))
print('Accuracy of CV - Decision Tree: %.4f (%.4f)' % (mean(DecTree_scores), stdev(DecTree_scores)))
print('Accuracy of CV - Random Forest: %.4f (%.4f)' % (mean(RFTree_scores), stdev(RFTree_scores)))
```



```
Accuracy of CV - KNN: 0.9560 (0.0021)
Accuracy of CV - Decision Tree: 0.9538 (0.0019)
Accuracy of CV - Random Forest: 0.9682 (0.0013)
```


Random Forest remains the best model overall in both F1-score and cross-validation accuracy.

```
## Plotting ROC Curve for best performing Model (RF Classifier)
```

```
from sklearn.metrics import roc_curve
```

```
fpr, tpr, thresholds = roc_curve(y_test, RFTree_pred)
```

```
plt.figure(figsize=(6,4))
plt.plot(fpr, tpr, linewidth=2)
plt.plot([0,1], [0,1], 'k--')
plt.rcParams['font.size'] = 12
plt.title('ROC curve for RF Diabetes')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.show()
```



```
from sklearn.metrics import roc_auc_score
```

```
ROC_AUC = roc_auc_score(y_test, RFTree_pred)
print('ROC AUC : {:.4f}'.format(ROC_AUC))
```



```
ROC AUC : 0.8481
```

```
Cross_validated_ROC_AUC = cross_val_score(RFTree, X_train, y_train, cv=10, scoring='roc_auc').mean()
```

```
print('Cross validated ROC AUC : {:.4f}'.format(Cross_validated_ROC_AUC))
```



```
Cross validated ROC AUC : 0.9580
```

It is evident that ROC AUC is better after the cross-validation. So, the cross-validation improved the model.

✓ Your model performs well overall (AUC = 0.8481, CV AUC = 0.9580).

✓ Cross-validation suggests the model generalizes well but may have some overfitting.

✓ Improving the recall for diabetics is crucial in medical applications (for example adjust thresholds, feature selection).

Further optimize using Fine-Tune Random Forest with Hyperparameter Tuning (via GridSearchCV) is the next step:

```
# Fine-Tuning Random Forest using GridSearchCV for maximizing accuracy & F1-score
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import GridSearchCV
```

```
# Define Random Forest model
rf = RandomForestClassifier(random_state=42)
```

```
# Define hyperparameters for tuning
param_grid_rf = {
    'n_estimators': [50, 100, 150], # Number of trees
    'max_depth': [10, 20, None], # Maximum depth of each tree
    'min_samples_split': [2, 5, 10] # Minimum samples to split an internal node
}
```

```
# Perform GridSearchCV
```

```

# Perform GridSearchCV
grid_rf = GridSearchCV(rf, param_grid_rf, cv=5, scoring='accuracy', n_jobs=-1)
grid_rf.fit(X_train, y_train)

# Best parameters and score
print("Best Parameters for Random Forest:", grid_rf.best_params_)
print("Best Accuracy for Random Forest:", grid_rf.best_score_)

↗ Best Parameters for Random Forest: {'max_depth': 10, 'min_samples_split': 2, 'n_estimators': 100}
Best Accuracy for Random Forest: 0.9718625

```

Increased accuracy from 96.82% → 97.19%

Two possible improvements:

Adjust the Decision Threshold (if recall is too low).

Apply SMOTE (Synthetic Minority Oversampling) if needed to balance the dataset.

```

# Step 1 Adjust Decision Threshold & Evaluate Model
classification models use 0.5 as the probability threshold
#If probability ≥ 0.5, the model predicts diabetes (1). If probability < 0.5, the model predicts no diabetes (0).

# Ensure 'best_rf' is assigned from GridSearchCV
best_rf = grid_rf.best_estimator_ # Extract the best Random Forest model

# Get predicted probabilities for diabetes (class 1)
y_probs = best_rf.predict_proba(X_test)[:, 1]

# Adjust the threshold from default 0.5 to a lower value (e.g., 0.35)
new_threshold = 0.35
y_pred_new = (y_probs >= new_threshold).astype(int)

# Evaluate new predictions
from sklearn.metrics import classification_report
print(f"Classification Report with Adjusted Threshold ({new_threshold}):\n", classification_report(y_test, y_pred_new))

```

```

↗ Classification Report with Adjusted Threshold (0.35):

```

	precision	recall	f1-score	support
0	0.97	1.00	0.99	18292
1	0.98	0.69	0.81	1708
accuracy			0.97	20000
macro avg	0.97	0.85	0.90	20000
weighted avg	0.97	0.97	0.97	20000

Diabetes recall improved from ~61% (before) to 69%.

If avoiding false positives (over-diagnosis) is more important → Use 0.35 (98% precision, but misses more diabetics).

```

#Step 2 Automatically Find the Best Threshold for Higher Recall

import numpy as np
from sklearn.metrics import precision_recall_curve, classification_report

# Compute Precision-Recall curve
precision, recall, thresholds = precision_recall_curve(y_test, y_probs)

# Exclude edge cases where threshold is too low or too high
valid_thresholds = thresholds[(recall[:-1] >= 0.75) & (recall[:-1] <= 0.85)] # Recall between 75-85%

if len(valid_thresholds) > 0:
    best_threshold = valid_thresholds[0] # Select the first reasonable threshold
else:
    best_threshold = 0.5 # Default to 0.5 if no valid threshold is found

print(f"Best Threshold for Balanced Recall & Precision: {best_threshold:.2f}")

# Apply the best threshold
y_pred_best = (y_probs >= best_threshold).astype(int)

# Evaluate new predictions with the best threshold
print(f"Classification Report with Best Threshold ({best_threshold}):\n",
      classification_report(y_test, y_pred_best, zero_division=1))

```



```

Best Threshold for Balanced Recall & Precision: 0.14
Classification Report with Best Threshold (0.13777394660575534):

```

	precision	recall	f1-score	support
0	0.99	0.95	0.97	18292
1	0.63	0.85	0.72	1708
accuracy			0.94	20000
macro avg	0.81	0.90	0.85	20000
weighted avg	0.96	0.94	0.95	20000

If detecting more diabetics (higher recall) is more important → Use 0.14 (85% recall, detects more actual diabetics but lower precision).

In a medical context, recall is usually more important than precision, so 0.14 may be better for diabetes screening.

✓ SMOTE

Since our dataset is imbalanced (fewer diabetic cases than non-diabetics), we can use SMOTE (Synthetic Minority Oversampling Technique) to generate synthetic diabetic cases and improve recall without losing accuracy.

```

from imblearn.over_sampling import SMOTE
from sklearn.model_selection import train_test_split
import numpy as np

# Apply SMOTE only on the training data
smote = SMOTE(sampling_strategy='auto', random_state=42) # auto balances minority class
X_train_sm, y_train_sm = smote.fit_resample(X_train, y_train)

# Check new class distribution after SMOTE
print("Class distribution before SMOTE:", np.bincount(y_train))
print("Class distribution after SMOTE:", np.bincount(y_train_sm))

Class distribution before SMOTE: [73208 6792]
Class distribution after SMOTE: [73208 73208]

from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import classification_report

# Use the fine-tuned best Random Forest model
rf_smote = RandomForestClassifier(max_depth=10, min_samples_split=2, n_estimators=100, random_state=42)

# Train on SMOTE-balanced dataset
rf_smote.fit(X_train_sm, y_train_sm)

# Get predictions on the test set
y_probs_sm = rf_smote.predict_proba(X_test)[: , 1] # Get probabilities for class 1 (diabetes)

# Apply the best decision threshold (from previous step, e.g., 0.14)
threshold = 0.14
y_pred_sm = (y_probs_sm >= threshold).astype(int)

# Evaluate model after SMOTE
print(f"Classification Report After SMOTE (Threshold {threshold}):\\n", classification_report(y_test, y_pred_sm))

```

```

Classification Report After SMOTE (Threshold 0.14):

```

	precision	recall	f1-score	support
0	0.98	0.88	0.93	18300
1	0.40	0.82	0.53	1700
accuracy			0.88	20000
macro avg	0.69	0.85	0.73	20000
weighted avg	0.93	0.88	0.90	20000

Recall for diabetes increased significantly from ~69% to 82%!

Precision for diabetes dropped to 40%, meaning more false positives (non-diabetics mistakenly classified as diabetics).

Overall accuracy remains strong (88%) despite recall improvements.

✓ Feature Engineering

```
# Split data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42, stratify=y)

# Apply SMOTE to balance the dataset
smote = SMOTE(random_state=42)
X_train_smote, y_train_smote = smote.fit_resample(X_train, y_train)

# Standardize features
scaler = StandardScaler()
X_train_smote = scaler.fit_transform(X_train_smote)
X_test = scaler.transform(X_test)

# ... (Your existing code for feature engineering) ...

from sklearn.decomposition import PCA
from sklearn.feature_selection import SelectKBest, f_classif
from sklearn.preprocessing import PolynomialFeatures

# Apply Polynomial Features
poly = PolynomialFeatures(degree=2, interaction_only=True, include_bias=False)
X_train_poly = poly.fit_transform(X_train_smote) # Use X_train_smote here
X_test_poly = poly.transform(X_test)

# Apply PCA for dimensionality reduction
pca = PCA(n_components=0.95) # Retain 95% variance
X_train_pca = pca.fit_transform(X_train_poly)
X_test_pca = pca.transform(X_test_poly)

# Apply Feature Selection
selector = SelectKBest(score_func=f_classif, k=10) # Keep top 10 features
X_train_selected = selector.fit_transform(X_train_pca, y_train_smote)
X_test_selected = selector.transform(X_test_pca)

print(f"Original Features: {X_train.shape[1]}, After Feature Engineering: {X_train_selected.shape[1]}")
```

🔄 Original Features: 6, After Feature Engineering: 10

Optimizing your models to see the impact of feature selection on performance.

```
#Retrain Random Forest with Feature-Selected Data
#Run this to train and evaluate Random Forest after feature selection:
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import recall_score, f1_score, roc_auc_score

# Train optimized Random Forest model
rf_optimized = RandomForestClassifier(n_estimators=300, max_depth=None, min_samples_split=5, class_weight="balanced", random_state=42)
rf_optimized.fit(X_train_selected, y_train_smote)

# Make predictions
y_pred_rf_optimized = rf_optimized.predict(X_test_selected)

# Evaluate performance
recall_rf_opt = recall_score(y_test, y_pred_rf_optimized)
f1_rf_opt = f1_score(y_test, y_pred_rf_optimized)
roc_auc_rf_opt = roc_auc_score(y_test, rf_optimized.predict_proba(X_test_selected)[:, 1])

print(f"Optimized Random Forest - Recall: {recall_rf_opt:.4f}, F1 Score: {f1_rf_opt:.4f}, ROC-AUC: {roc_auc_rf_opt:.4f}")
```

🔄 Optimized Random Forest - Recall: 0.8071, F1 Score: 0.6469, ROC-AUC: 0.9638

ROC-AUC (96.38%) – Excellent Performance. The ROC-AUC score of 0.9638 suggests the model has outstanding discrimination ability between diabetics and non-diabetics.

The F1-score is slightly lower than the earlier result (because improving recall often lowers precision).

Next Steps: Fine-Tune the Threshold Again – A slight increase to 0.15 - 0.20

✓ # Summary of codes use to date

Step 1. Splits into train-test sets: Uses an 80-20 split while maintaining class proportions.

Step 2. Applies SMOTE: Balances the dataset by oversampling the minority class.

Step 3. Standardizes the features: Ensures numeric consistency.

Step 4. Trains and evaluates models: Uses Logistic Regression, Random Forest, Gradient Boosting Classifier, Support Vector Machine (SVM), and (to do - XGBoost with class weighting.)

Step 5: Evaluates the model using the following Metrics:

Recall: Focuses on correctly identifying positive cases. F1 Score: Harmonic mean of precision & recall. ROC-AUC Score: Measures true positive vs. false positive rates.

Step 6. Uses GridSearchCV to find the best hyperparameters.

Step 7. More Fine-Tuning:

Now tunes Logistic Regression, Random Forest, and XGBoost for best performance.

Step 8 Displays best parameters and final metrics after hyperparameter tuning.

```
import time

def train_and_evaluate(model, X_train, y_train, X_test, y_test):
    start_time = time.time()
    model.fit(X_train, y_train)
    train_time = time.time() - start_time

    start_time = time.time()
    y_pred = model.predict(X_test)
    pred_time = time.time() - start_time

    recall = recall_score(y_test, y_pred)
    f1 = f1_score(y_test, y_pred)

    print(f"{model.__class__.__name__} - Train Time: {train_time:.4f}s, Prediction Time: {pred_time:.4f}s, Recall: {recall:.4f}, F1: {f1:.4f}")

# Load the New Dataset #2
# dataset 2

import pandas as pd # Use pd instead of pd2

# Correct RAW URL
url = "https://raw.githubusercontent.com/AsmaShaikhTMU/Projects/main/diabetes_012_health_indicators_BRFSS2015.csv"

# Read CSV file
df2 = pd.read_csv(url)

# Display first few rows
df2.head()
```

5 rows × 22 columns

	Diabetes_012	HighBP	HighChol	CholCheck	BMI	Smoker	Stroke	HeartDiseaseorAttack	PhysActivity	Fruits	...	AnyHealthcare	NoDoct
0	0.0	1.0	1.0	1.0	40.0	1.0	0.0	0.0	0.0	0.0	...	1.0	
1	0.0	0.0	0.0	0.0	25.0	1.0	0.0	0.0	1.0	0.0	...	0.0	
2	0.0	1.0	1.0	1.0	28.0	0.0	0.0	0.0	0.0	1.0	...	1.0	
3	0.0	1.0	0.0	1.0	27.0	0.0	0.0	0.0	1.0	1.0	...	1.0	
4	0.0	1.0	1.0	1.0	24.0	0.0	0.0	0.0	1.0	1.0	...	1.0	

```
print(df2.head())
```

```

Diabetes_012  HighBP  HighChol  CholCheck  BMI  Smoker  Stroke  \
0      0.0      1.0      1.0      1.0  40.0      1.0      0.0
1      0.0      0.0      0.0      0.0  25.0      1.0      0.0
2      0.0      1.0      1.0      1.0  28.0      0.0      0.0
3      0.0      1.0      0.0      1.0  27.0      0.0      0.0
4      0.0      1.0      1.0      1.0  24.0      0.0      0.0

HeartDiseaseorAttack  PhysActivity  Fruits  ...  AnyHealthcare  \
0      0.0      0.0      0.0  ...      1.0
1      0.0      1.0      0.0  ...      0.0
2      0.0      0.0      1.0  ...      1.0
3      0.0      1.0      1.0  ...      1.0
4      0.0      1.0      1.0  ...      1.0

NoDocbcCost  GenHlth  MentHlth  PhysHlth  DiffWalk  Sex  Age  Education  \
0      0.0      5.0      18.0      15.0      1.0  0.0  9.0      4.0
1      1.0      3.0      0.0      0.0      0.0  0.0  7.0      6.0
2      1.0      5.0      30.0      30.0      1.0  0.0  9.0      4.0
3      0.0      2.0      0.0      0.0      0.0  0.0  11.0     3.0
4      0.0      2.0      3.0      0.0      0.0  0.0  11.0     5.0

Income
0      3.0
1      1.0
2      8.0
3      6.0
4      4.0

```

[5 rows x 22 columns]

```

# Compare feature distributions
for col in X_train.columns:
    plt.figure(figsize=(6,4))
    sns.kdeplot(X_train[col], label="Training Data", fill=True)
    sns.kdeplot(df2[col], label="Dataset #2", fill=True)
    plt.title(f"Feature Distribution: {col}")
    plt.legend()
    plt.show()

```

